1. Select the option that correctly completes the sentence:

1 point

Training a model using labeled data and using this model to predict the labels for new data is known as _____.

○ Unsupervised Learning

○ Density Estimation

○ Clustering

◉ Supervised Learning

2. Select the option that correctly completes the sentence:

1 point

Modeling the features of an unlabeled dataset to find hidden structure is known as _____.

○ Classification

○ Supervised Learning

○ Regression

◉ Unsupervised Learning

3. Select the option that correctly completes the sentence:

1 point

Training a model using categorically labelled data to predict labels for new data is known as _____.

○ Regression

◉ Classification

○ Feature Extraction

○ Clustering

**4.** Select the option that correctly completes the sentence:

1 point

Training a model using labelled data where the labels are continuous quantities to predict labels for new data is known as _____.

○ Feature Extraction

◉ Regression

○ Classification

○ Clustering

5. Using the data for classes 0, 1, and 2 plotted below, what class would a KNeighborsClassifier classify the new point as for k = 1 and k = 3?

1 point

- ○   • k=1: Class 0
  - • k=3: Class 1

- ○   • k=1: Class 1
  - • k=3: Class 0

- ○   • k=1: Class 2
  - • k=3: Class 1

- ○   • k=1: Class 0
  - • k=3: Class 2

- ◉   • k=1: Class 1
  - • k=3: Class 2

6. Which of the following would *not* be something you specify for a nearest neighbor classifier algorithm?   `1 point`

- ☑ The number of initial clusters to create from the training set
- ☐ An optional weighting function on the neighbor points
- ☐ How many nearest neighbors to examine
- ☐ A distance metric that finds neighbor points in the training set
- ☐ A method for pooling the classes of neighbor points to make a final classification decision

7. Why is it important to examine your dataset as a first step in applying machine learning? (Select all that apply):   `1 point`

- ☑ See what type of cleaning or preprocessing still needs to be done
- ☑ You might notice missing data
- ☑ Gain insight on what machine learning model might be appropriate, if any
- ☑ Get a sense for how difficult the problem might be
- ☐ It is not important

8. The key purpose of splitting the dataset into training and test sets is: `1 point`

- ○ To speed up the training process
- ◉ To estimate how well the learned model will generalize to new data
- ○ To reduce the number of features we need to consider as input to the learning algorithm
- ○ To reduce the amount of labelled data needed for evaluating classifier accuracy

9. The purpose of setting the random_state parameter in train_test_split is: (Select all that apply)　　`1 point`

- ☑ To make experiments easily reproducible by always using the same partitioning of the data
- ☐ To split the data into similar subsets so that bias is not introduced into the final results
- ☐ To avoid predictable splitting of the data
- ☐ To avoid bias in data splitting

10. Given a dataset with 10,000 observations and 50 features plus one label, what would be the dimensions of X_train, y_train, X_test, and y_test? Assume a train/test split of 75%/25%.　　`1 point`

- ○
  - X_train: (10000, 28)
  - y_train: (10000,)
  - X_test: (10000, 12)
  - y_test: (10000,)

- ○
  - X_train: (2500, 50)
  - y_train: (2500,)
  - X_test: (7500, 50)
  - y_test: (7500,)

- ○
  - X_train: (2500,)
  - y_train: (2500, 50)
  - X_test: (7500,)
  - y_test: (7500, 50)

- ◉
  - X_train: (7500, 50)
  - y_train: (7500,)
  - X_test: (2500, 50)
  - y_test: (2500,)

- ○
  - X_train: (10000, 50)
  - y_train: (10000,)
  - X_test: (10000, 50)
  - y_test: (10000,)

11. Which of the following statements about the effect of the number of neighbors parameter k is true?　　`1 point`

- ☑ For larger values of k (e.g. k = 10), the regions in feature space assigned to different classes generally become less fragmented and have smoother decision boundaries between the regions, compared to results using smaller values of k.
- ☐ For small values of k (e.g. k = 1), the classifier will tend to be much <u>less</u> sensitive to noise, mislabeled data, and other sources of variation for individual data points.
- ☑ For small values of k (e.g. k = 1), the classifier will tend to be much <u>more</u> sensitive to noise, mislabeled data, and other sources of variation for individual data points.
- ☐ Nearest neighbor classifier accuracy is never significantly affected by the choice of k.